

Akshansh Vij
102317156
2Q16

Lab Assignment-11

Cognitive Computing UCS420 Landing AI (Vision-based Cognitive System)

This project is a classification model built using LandingAI's computer vision platform to distinguish between two types of stationary items: pens and scales. Utilizing the **Classification** project type, the model applies a multi-class classification approach to categorize input images into one of two classes: "**Keyboard(B)**" or "**Mouse(A)**". By training on labeled images of both categories, the model learns to accurately identify and differentiate between these common stationary items based on visual features.

Project Creation Page

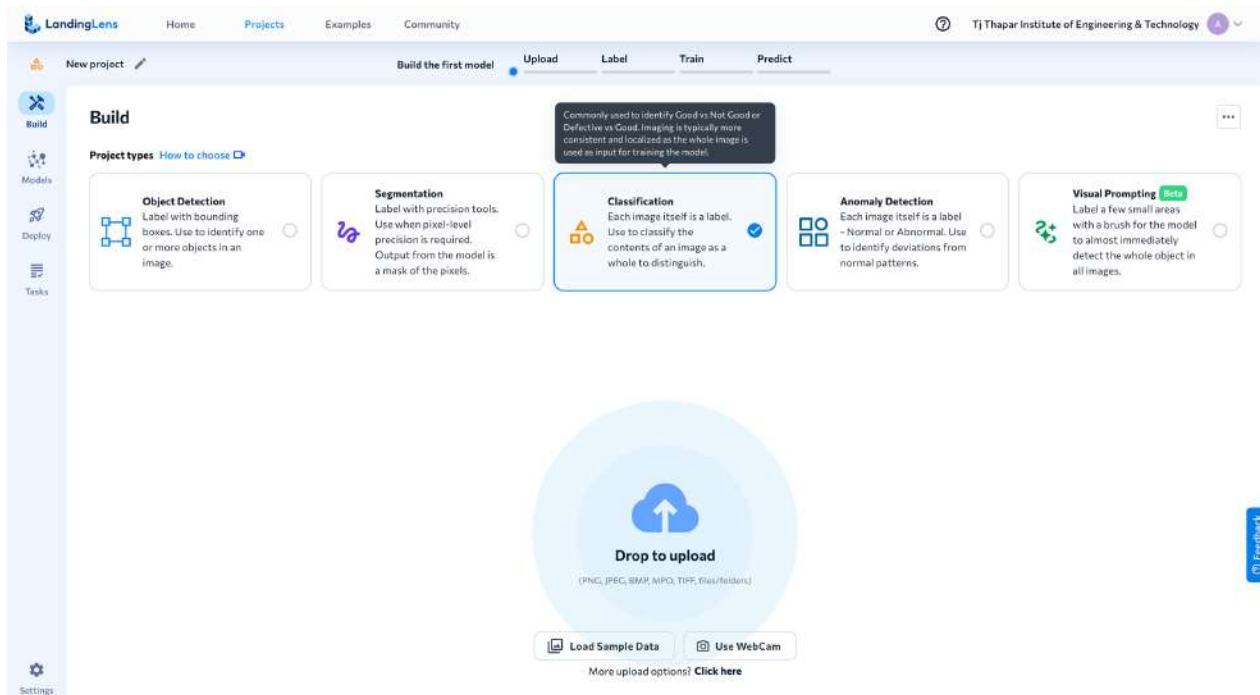


Image upload



Classified images upload

Drag and drop folders/images here, or click to select files

Split: unassigned

Metadata: Add metadata

Tag: Add Tag

Upload 18 Image(s)

18 Image(s) Ready for Upload

Upload Preview

Class A

10 Image(s)



Class B

8 Image(s)



Training Process

The screenshot shows the LandingLens 'Build' stage. The top navigation bar includes 'Home', 'Projects', 'Examples', and 'Community'. The user is logged in as 'TJ Thapar Institute of Engineering & Technology'. The main workflow is 'Build the first model', with steps 'Upload', 'Label', 'Train', and 'Predict'. The 'Train' step is currently active. On the left, a sidebar shows 'Build', 'Models', 'Deploy', and 'Tasks'. The main area displays '18 images' in a grid. The first row shows five images of keyboards, each labeled 'Class B'. The second row shows three more keyboard images labeled 'Class B', followed by two images of a computer mouse labeled 'Class A'. A 'Train' button is visible in the top right of the image grid. On the right, a panel for 'Model-04-24-2025_1' shows the training progress: 'Preparing data snapshot' (checked), 'Provisioning GPU' (checked), 'Training & learning' (in progress), and 'Calculating performance' (pending). An 'End Training Now' button is at the bottom of this panel.

This screenshot shows the same LandingLens interface as the first, but with a modal window open for 'Model-04-24-2025_1'. The modal provides a detailed view of the training process: 1. 'Preparing the snapshot of your data' (checked) with the note 'snapshot AutoGenerated-04-24-2025_1 with 18 images.' 2. 'Provisioning GPU' (checked) with the note 'Warming up a virtual computer in the cloud.' 3. 'Training & learning' (in progress). 4. 'Calculating model performance' (pending). A 'Show Training Configuration' link is at the bottom. The background shows the same image grid, but it is dimmed. The right sidebar panel remains visible.

Testing Results

The image displays two screenshots of the LandingLens application interface, showing the results of a model's predictions on different input images.

Top Screenshot:

- The interface shows the "Try this model" window.
- The input image is a photograph of a black computer keyboard.
- The prediction result is displayed as "Class B 0.98".
- A "Deploy" button is visible in the top right corner of the prediction window.

Bottom Screenshot:

- The interface shows the "Try this model" window.
- The input image is a photograph of a person's hand holding a black computer mouse.
- The prediction result is displayed as "Class A 0.85".
- A "Deploy" button is visible in the top right corner of the prediction window.

The background of both screenshots shows the LandingLens application interface, including the "Build" tab, "Filter" button, and "Model list" section.